

EECS-3421N: TEST #2

Electrical Engineering & Computer Science
Lassonde School of Engineering
York University

Family Name: _____
Given Name: _____
Student#: _____
EECS Account: _____

Instructor: Parke Godfrey

Exam Duration: 75 minutes

Term: Winter 2019

Instructions

- **rules**
 - The test is closed-note, closed-book. Use of a calculator is permitted.
- **answers**
 - Should you feel a question needs an assumption to be able to answer it, write the assumptions you need along with your answer.
 - If you need more room to write an answer, indicate where you are continuing the answer.
 - For multiple choice questions, choose *one* best answer for each of the following. There is no negative penalty for a wrong answer.
- **notation**
 - For schema, the underlined attributes indicate a table's primary key (and are, hence, not nullable). Attributes in *italics* are not nullable. Foreign keys are indicated by FK.
 - Assume *set* semantics for relational-algebra expressions.
- **points**
 - The number of points a given question is worth is marked.
 - There are five major parts worth 10 points each, for 50 points in total.

MARKING BOX	
1.	/10
2.	/10
3.	/10
4.	/10
5.	/10
Total	/50

HAPPY PI DAY!

1. [10pt] **General.** *Much choice!*

MULTIPLE CHOICE

Choose *one* best answer for each of the following. Each is worth one point. There is no negative penalty for a wrong answer.

In the *rare* case that you feel a clarification to your answer is needed, write a brief clarification on the side.

Let $|\mathbf{T}|$ denote the number of tuples in \mathbf{T} .

- (a) [1pt] SQL is derived from
- A. the *relational algebra*.
 - B. the *domain relational calculus*.
 - C. the *tuple relational calculus*.
 - D. XQuery for XML.
 - E. *SQL is a programming language, not a query language. It is not derived from any of the above.*
-

- (b) [1pt] The *relational algebra* is an algebra because
- A. it contains operators for arithmetic.
 - B. its operators map from the domain of operands back into the domain.
 - C. it allows one to order freely its operators to form expressions.
 - D. it contains a fixed, finite number of pre-defined operators.
 - E. *It is not really an algebra. But calling it an algebra sounds cool!*
-

- (c) [1pt] In a relational database system, if you join (natural join) tables \mathbf{R} and \mathbf{S} , but \mathbf{R} is empty (that is, it has no tuples),
- A. the system reports an error.
 - B. the answer set is an empty table.
 - C. the answer set is the same as table \mathbf{S} .
 - D. the answer set consists of just *one* row.
 - E. an answer set is returned; however, the results are system dependent.
-

- (d) [1pt] $\mathbf{R} \cap \mathbf{S}$ is equivalent to
- A. $\mathbf{R} - (\mathbf{R} - \mathbf{S})$
 - B. $\mathbf{R} - (\mathbf{S} - \mathbf{R})$
 - C. $(\mathbf{R} - \mathbf{S}) - \mathbf{R}$
 - D. $- ((-\mathbf{R}) \cup (-\mathbf{S}))$
 - E. *There is not enough information to answer this.*
-

- (e) What does the query “`select max(R.B) from R;`” return if \mathbf{R} is empty?
- A. An empty table of one column.
 - B. A table of one column with one row with the value $\langle \text{NULL} \rangle$.
 - C. A table of one column with one row with the value $\langle \text{INF} \rangle$.
 - D. An error message.
 - E. *Not enough information to determine.*
-

- (f) [1pt] In *relational algebra*, we can write any expression using *join* (\bowtie) using the following instead.
- A. π, \cap
 - B. π, \cup
 - C. π, σ, \times
 - D. $\pi, -$
 - E. *Join cannot be rewritten with any combination of the other R.A. operators.*
-

- (g) [1pt] Which of the following may evaluate to a different answer than $\mathbf{R} \bowtie (\mathbf{S} \bowtie \mathbf{T})$?
- A. $(\mathbf{R} \bowtie \mathbf{S}) \bowtie \mathbf{T}$
 - B. $\mathbf{R} \bowtie (\mathbf{T} \bowtie \mathbf{R})$
 - C. $(\mathbf{S} \bowtie \mathbf{T}) \bowtie \mathbf{R}$
 - D. $\mathbf{T} \bowtie (\mathbf{S} \bowtie \mathbf{R})$
 - E. *They all evaluate the same.*
-

- (h) [1pt] Consider the schema

$\mathbf{R}(\underline{\mathbf{A}}, \mathbf{B})$ FK (B) refs $\mathbf{R}(\mathbf{A})$

What is the *largest* that $|\mathbf{R} \bowtie \pi_{\mathbf{A} \rightarrow \mathbf{B}, \mathbf{B} \rightarrow \mathbf{A}}(\mathbf{R})|$ can be?

- A. 0
 - B. $|\mathbf{R}|$
 - C. $\frac{1}{2}|\mathbf{R}|$
 - D. $2|\mathbf{R}|$
 - E. $|\mathbf{R}|^2$
-

- (i) [1pt] Consider table $\mathbf{R}(\underline{\mathbf{A}}, \mathbf{B})$ for which B is of type *integer* and $|\mathbf{R}| = n > 0$. How many tuples will the query

`select A from R where B <= 13 or B > 13;`

return?

- A. 0
 - B. $\frac{1}{2}n$
 - C. n
 - D. n^2
 - E. *There is not enough information to answer this.*
-

- (j) [1pt] Consider the relations $\mathbf{R}(\underline{\mathbf{A}}, \underline{\mathbf{B}})$, $\mathbf{S}(\underline{\mathbf{B}}, \underline{\mathbf{C}})$, and $\mathbf{T}(\underline{\mathbf{C}}, \underline{\mathbf{A}})$.

One of these is not like the others. That is, one can evaluate differently than the other four. Which one?

- A. $\pi_{\mathbf{A}, \mathbf{B}}((\mathbf{R} \bowtie \mathbf{S}) \bowtie \mathbf{T})$
 - B. $\pi_{\mathbf{A}, \mathbf{B}}(\mathbf{R} \bowtie (\mathbf{T} \bowtie \mathbf{S}))$
 - C. $\mathbf{R} \bowtie \pi_{\mathbf{A}, \mathbf{B}}(\mathbf{S} \bowtie \mathbf{T})$
 - D. $\pi_{\mathbf{A}, \mathbf{B}}(\mathbf{R} \bowtie \mathbf{T}) \bowtie \pi_{\mathbf{A}, \mathbf{B}}(\mathbf{R} \bowtie \mathbf{S})$
 - E. $\pi_{\mathbf{A}, \mathbf{B}}(\mathbf{R} \bowtie \mathbf{S}) \bowtie \pi_{\mathbf{A}, \mathbf{B}}(\mathbf{S} \bowtie \mathbf{T})$
-

2. [10pt] **Relational Algebra.** *Riddle me this!*

SHORT ANSWER

For Questions 2a to 2d, use the *Colours* schema in Figure 1 on page 13 (as used in examples in class) to write *relational-algebra* queries for the English questions posed and vice versa.

(a) [2pt] Show products by `prod#` and `pname` that come in colour *orange*.

(b) [2pt] Show products by `prod#` and `pname` that are owned by at least two customers.

(c) [2pt] Show pairs of customers—*first* (`cust#`) and *fname* (`cname`) for the first customer and *second* (`cust#`) and *sname* (`cname`) for the second customer—such that the second customer owns an item in the first customer's favourite colour.

- (d) [2pt] Consider the following R.A. expression.

$$\pi_{\text{colour}}(\pi_{\text{colour,prod\#,name}}(\text{Product} \bowtie \text{Avail_colour}) \\ - \\ \pi_{\text{colour,prod\#,name}}(\sigma_{\text{cost} > \text{cost2}}(\\ \pi_{\text{colour,prod\#,name,cost}}(\text{Product} \bowtie \text{Avail_colour}) \\ \bowtie \\ \pi_{\text{colour,prod\#,name,cost} \rightarrow \text{cost2}}(\text{Product} \bowtie \text{Avail_colour}) \\)) \\)$$

State what the query asks in English.

Note that you will get *zero* credit if you use database terms in your answer! (E.g., “Well, the query first *joins* two tables, taking the *projection* of...” does not count!)

- (e) [2pt] Consider the following SQL query.

```
select distinct Z.A, X.C
from R Z, S X
where Z.B not in (
    select Y.B
    from S Y
    where X.C = Y.C);
```

Write an equivalent R.A. expression for it.

3. (10 points) **Queries in SQL.** *Ask me anything.*

EXERCISE

Consider the *Movie* schema in Figure 2 on page 13 for Questions 3a to 3d.

(a) [3pt] Write an SQL query that answers the following.

Report *directors* by **p#** and **name** with the movies that they have directed by **title**, **studio**, **year**, and **genre**.

Order by name, p#, year, studio, title (all ascending).

(b) [2pt] Write an SQL query that answers the following.

Report *directors* by **p#** and **name** with the number of movies that they have directed as **#movies**. You may assume that every director has directed some movie.

Order by name, p# (both ascending).

(c) [3pt] Consider the following SQL query.

```
select  P.p#, P.name,
        ( select count(*)
          from Cast C
         where P.p# = C.p#
        ) as #roles
from Person P;
```

Assume that everyone who has been cast in a movie is an *actor*, and that every *actor* has been in some movie.

Write an SQL query that means the same thing, but that uses *no* nested (sub-) queries.

(d) [2pt] Consider the following SQL query.

```
select P.name, P.gender, C.role, C.minutes, M.title, M.studio, M.year
from Person P, Cast C, Movie M, Authored A, Person W
where C.actor = P.p#
   and C.title = M.title and C.studio = M.studio and C.year = M.year
   and M.genre = 'SciFi'
   and A.title = M.title and A.year = M.year
   and A.writer = W.p#
   and W.name = 'Hampton Fancher';
```

State what the query asks in English.

Note that you will get *zero* credit if you use database terms in your answer! (E.g., “Well, the query first *joins* two tables, taking the *projection* of...” does not count!)

4. (10 points) **Query Logic.** *Fascinating.*ANALYSIS

- (a) [2pt] Consider the schema

 $\mathbf{S}(\underline{B}, C)$ $\mathbf{R}(\underline{A}, B)$ FK (B) refs $\mathbf{S}(B)$ What is $|\mathbf{R} \bowtie \mathbf{S}|$? (That is, what is the cardinality of $\mathbf{R} \bowtie \mathbf{S}$?)

Give your answer in the simplest form.

-
- (b) [2pt] You execute the following command against the database system for the
- Colours*
- database (schema in Figure 1 on page 13).

```
insert into Item (item#, prod#, cust#, colour, date_sold)
values
    (1729, 23, 13, 'hot pink', '03/14/2019');
```

Is this tuple now guaranteed to have been added to the table **Item**?

Why or why not?

-
- (c) [2pt] Say that “
- $\mathbf{R}.A <> \mathbf{S}.B$
- ” appears as a predicate in the
- where*
- clause of an SQL query. To what does the predicate evaluate when
- $\mathbf{R}.A$
- is
- null*
- and
- $\mathbf{S}.B$
- is
- null*
- ?

- (d) [2pt] Consider the *Colours* database (schema in Figure 1 on page 13) and the following SQL query.

```
select distinct C.cust#, P.prod#
from Customer C, Avail_Colours A
where C.fav_colour <> A.colour;
```

State what the query asks in English.

Note that you will get *zero* credit if you use database terms in your answer! (E.g., “Well, the query first *joins* two tables, taking the *projection* of...” does not count!)

-
- (e) [2pt] Consider the following SQL query.

```
select A, B
from (values
      (3.1416, 'x'),
      (2.7183, 'y')
    ) as R (A, B)
except
select C
from (values
      (3.1416)
    ) as S (C);
```

How does the query evaluate?

5. [10pt] **The SQL Language.** *Speaking quite logically.*

SHORT ANSWER

Consider the following for Questions 5a, 5b, and 5c.

Dr. Mark Dogfury has written his own **SQL** database system that he calls **WoofSQL**. Unfortunately, he forgot to implement “distinct”, and he forgot to implement “having”! However, he did implement the rest of **SQL**.

- (a) [3pt] Can you still write queries in **WoofSQL** that eliminate duplicate values, despite missing “distinct”?

Why or why not?

-
- (b) [2pt] Can you still write all the queries that you can in regular, *full* **SQL** in **WoofSQL**, even though you do not have the *having* clause?

Why or why not?

- (c) [2pt] Dr. Dogfury implemented *distinct* and *having* for **WoofSQL v2**, but somehow the *intersect* operator got left out of **v2**!

Can you still write all the queries that you can in regular, *full* SQL in **WoofSQL v2**, even though you do not have the *intersect* operator?

Why or why not?

-
- (d) [3pt] Say that you know that the **Product** table in the *Colours* database (schema in Figure 1 on page 13) has 45 tuples in it.

How many tuples are returned by the following query?

```
select prod#, pname
from Product as P,
     (values ('A'), ('B'), ('C'))
  as Q (X);
```

EXTRA SPACE

RELAX. TURN IN YOUR TEST. RETURN TO THE WILD! EAT SOME PIE!

REFERENCE

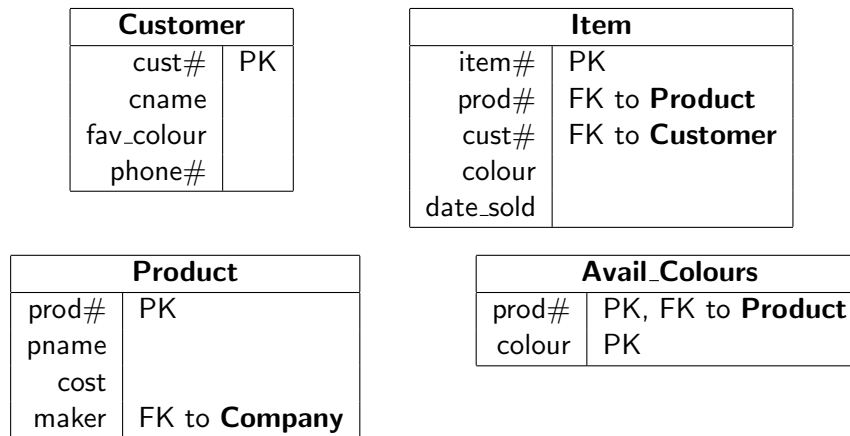
*(Detach this page for convenience, if you want.)***Schema for the Colours Database.**

Figure 1: Colours Schema.

Schema for the Movie Database.

Person(p#, name, birthdate, nationality, gender)
Actor(p#, aguild#)
 FK (p#) refs **Person**
Director(p#, dguild#)
 FK (p#) refs **Person**
Writer(p#, wguild#)
 FK (p#) refs **Person**
Studio(name)
ScreenPlay(title, year)
Authored(title, year, writer)
 FK (title, year) refs **ScreenPlay**
 FK (writer) refs **Writer** (p#)
Movie(title, studio, year, genre, director, length)
 FK (studio) refs **Studio** (name)
 FK (title, year) refs **ScreenPlay**
 FK (director) refs **Director** (p#)
Cast(title, studio, year, role, actor, minutes)
 FK (title, studio, year) refs **Movie**
 FK (actor) refs **Actor** (p#)
Affiliated(director, studio)
 FK (director) refs **Director** (p#)
 FK (studio) refs **Studio** (name)

Figure 2: Movie Schema.

REFERENCE

The Relational-Algebra Operators.

σ : selection
 π : projection
 \bowtie : join
 \times : cross product
 \cup : union
 \cap : intersection
 $-$: difference
 ρ : rename