

# EECS-3421A: TEST #2

## “Queries”

*Electrical Engineering & Computer Science*

*Lassonde School of Engineering*

**York University**

**Family Name:** \_\_\_\_\_

**Given Name:** \_\_\_\_\_

**Student#:** \_\_\_\_\_

**EECS Account:** \_\_\_\_\_

**Instructor:** Parke Godfrey

**Exam Duration:** 75 minutes

**Term:** Fall 2016

### Instructions

- **rules**
  - The test is closed-note, closed-book. Use of a calculator is permitted.
- **answers**
  - Should you feel a question needs an assumption to be able to answer it, write the assumptions you need along with your answer.
  - If you need more room to write an answer, indicate where you are continuing the answer.
  - For multiple choice questions, choose *one* best answer for each of the following. There is no negative penalty for a wrong answer.
- **notation**
  - For schema, the underlined attributes indicate a table’s primary key (and are, hence, not nullable). Attributes in *italics* are not nullable. Foreign keys are indicated by FK.
  - Assume *set* semantics for relational-algebra expressions.
- **points**
  - The number of points a given question is worth is marked.
  - There are five major parts worth 10 points each, for 50 points in total.

MARKING BOX	
<b>1.</b>	/10
<b>2.</b>	/10
<b>3.</b>	/10
<b>4.</b>	/10
<b>5.</b>	/10
<b>Total</b>	/50

1. [10pt] **Relational Algebra.** *Quadratic queries!*

[SHORT ANSWER]

For Questions 1a to 1c, use the Colours schema in Figure 1 on page 15 (as used in examples in class) to write *relational-algebra* queries for the English questions posed.

---

- a. [2pt] Show customers by `cust#` and `cname` who have bought a *pink* (colour) *Lamborghini* (`prod#`).
- 

- b. [2pt] Show products by `prod#` and `pname` that are owned by at least two customers.
- 

- c. [2pt] For each available colour (`colour`), show the most expensive product by `prod#` and `pname` that comes in that colour. (In case of ties for most expensive, list all in the tie.)

<b>R</b>	
A	B
a	b
c	b
e	f
g	h
i	h

<b>S</b>	
B	C
f	b
b	d
h	a
h	c
b	e

<b>T</b>	
A	C
g	a
a	b
i	c
e	d
c	e

Consider the three tables **R**, **S**, & **T** above for Questions 1d & 1e.

---

d. [2pt] Show the results of  $\mathbf{R} \bowtie \mathbf{S}$ .

---

e. [2pt] Show the results of  $\mathbf{R} \bowtie (\mathbf{S} \bowtie \mathbf{T})$ .

2. (10 points) **SQL.** *Some Quidditch League!*

[EXERCISE]

Consider the Movie database with the schema in Figure 2 on page 15 for the questions below.

---

- a. [5pt] Write an SQL query for the following.

List each actor by name, gender, role, and the minutes they appear on screen in that role in a given movie by title, studio, and year such that “Hampton Fancher” was an author of the screenplay of the movie and the movie’s genre is “SciFi”.

- b. [5pt] For each movie, report how many male actors and how many female actresses—`gender = 'M'` for *male* and `gender = 'F'` for *female*—were *cast* in the movie in columns `male` and `female`, respectively.

Only count a given actor or actress *once* per movie; that is, that a person may have played several *characters* (*roles*) in the movie does *not* count multiple times.

For full credit, the column `male` or `female` should report '0' (zero) if there were no actors or actresses, respectively, in the movie.

---

3. (10 points) **Query Logic.** Take the *L* train to *Q* Street.

[ANALYSIS]

- 
- a. [2pt] State in *plain, concise English* what the following SQL query over the *Movie* database does. (See the Movie schema in Figure 2 on page 15.)

Note that you will get *zero* credit if you use database terms in your answer! (E.g., “Well, the query first *joins* two tables, taking the *projection* of...” does not count!)

```
select P.p#, P.name as actor, D.p# as d#, D.name as director
from Person P, Person D
where not exists (
    select *
    from Movie M
    where M.director = D.p#
    and not exists (
        select *
        from Cast C
        where C.actor = P.p#
        and C.title = M.title
        and C.studio = M.studio
        and C.year = M.year
    )
)
and exists (
    select *
    from Movie M
    where M.director = D.p#
);
```

- 
- b. [2pt] Is the *having* clause *logically redundant* in SQL? That is, could one always write the “same” query not using the *having* clause?

Explain briefly why or why not.

- c. [2pt] Given  $\mathbf{R}(\underline{\mathbf{A}}, \underline{\mathbf{B}})$  and  $\mathbf{S}(\underline{\mathbf{B}}, \underline{\mathbf{C}})$ , rewrite

$$(\pi_{\mathbf{A}}(\mathbf{R}) \times \pi_{\mathbf{C}}(\mathbf{S})) - \pi_{\mathbf{A}, \mathbf{C}}(\mathbf{R} \bowtie \mathbf{S})$$

as an equivalent SQL query.

- 
- d. [2pt] Consider relations  $\mathbf{R}(\underline{\mathbf{A}}, \underline{\mathbf{B}})$ ,  $\mathbf{S}(\underline{\mathbf{B}}, \underline{\mathbf{C}})$ , and  $\mathbf{T}(\underline{\mathbf{C}}, \underline{\mathbf{D}})$ .

Can  $((\mathbf{R} \bowtie \mathbf{S}) \bowtie \mathbf{T})$  and  $((\mathbf{R} \bowtie \mathbf{T}) \bowtie \mathbf{S})$  evaluate to different answer sets, or must they evaluate to the same answer set?

*Show* how they could evaluate differently, *or argue* why they must evaluate to the same.

- 
- e. [2pt] Consider the relations  $\mathbf{R}(\underline{\mathbf{A}}, \mathbf{B})$  and  $\mathbf{S}(\underline{\mathbf{C}}, \mathbf{D})$ , and query

```
select distinct A, B from R
where exists (select * from S where A = C);
```

Write this as a relational-algebra expression.

---

4. [10pt] **Normalization.** *It's the end of the world as we know it...*

[ANALYSIS]

---

- a. [3pt] Consider the relation **T** with attributes A, B, C, D, and E, and with the following functional dependencies (FDs):

$$\begin{array}{ll} AB \mapsto C & C \mapsto A \\ AD \mapsto E & C \mapsto D \end{array}$$

Dr. Datta Bas claims that the decomposition step of **T** into ABC & CDE is a *lossless-join* decomposition step.

Explain convincingly *either* that he is correct *or* that he is wrong.

- 
- b. [2pt] Consider the relation **R** with attributes A, B, and C, and with just the one following functional dependency (FD):

$$A \mapsto BC$$

Consider the decomposition of **R** into AB and BC. Construct a counterexample (example tuples) that demonstrate that this decomposition is *not* lossless (that is, it is *lossy*).



For Questions 4c to 4e, consider the relation **R** with attributes A, B, C, and D, and with the following functional dependencies (FDs):

$$\begin{array}{ll} AB \mapsto C & A \mapsto D \\ BD \mapsto C & \end{array}$$

---

c. [1pt] What are the keys of **R**?

---

d. [3pt] Decompose **R** losslessly into a BCNF decomposition.

---

e. [1pt] Is your decomposition in your answer to Question 4d dependency preserving?  
Why or why not?

5. [10pt] **General.** ... and I feel fine. [MULTIPLE CHOICE]

Choose *one* best answer for each of the following. Each is worth one point. There is no negative penalty for a wrong answer.

In the *rare* case that you feel a clarification to your answer is needed, write a brief clarification on the side.

Let  $|\mathbf{T}|$  denote the number of tuples in  $\mathbf{T}$ .

- a. [1pt] Consider the schema

$\mathbf{R}(\underline{A}, B)$  FK (B) refs  $\mathbf{S}$

$\mathbf{S}(A, \underline{B})$  FK (A) refs  $\mathbf{R}$

Note that none of the attributes are nullable. Which of the following is guaranteed to produce as many as, or more, tuples than each of the others?

- A.  $\mathbf{R} \bowtie \mathbf{S}$
- B.  $\pi_A(\mathbf{R}) \bowtie \mathbf{S}$
- C.  $\pi_A(\mathbf{R}) \bowtie \pi_B(\mathbf{S})$
- D.  $\mathbf{R} \bowtie \pi_B(\mathbf{S})$
- E. *There is not enough information to answer this.*

- b. [1pt] Assume  $|\mathbf{R}| > 0$  and  $|\mathbf{S}| > 0$ . If one natural joins tables  $\mathbf{R}$  and  $\mathbf{S}$ , but  $\mathbf{R}$  and  $\mathbf{S}$  have no column names in common, then

- A. it is an *error*.
- B. the answer set is an empty table.
- C. it is an outer join.
- D. it is the same as  $\mathbf{R} \cap \mathbf{S}$ .
- E. it is the same as  $\mathbf{R} \times \mathbf{S}$ .

For Questions 5c & 5d, consider the schema

$\mathbf{R}(\underline{A}, B)$  FK (B) refs  $\mathbf{R}$  (A)

- c. [1pt] What is the *smallest* that  $|\mathbf{R} \bowtie \pi_{A \rightarrow B, B \rightarrow A}(\mathbf{R})|$  can be?

- A. 0
- B.  $|\mathbf{R}|$
- C.  $\frac{1}{2}|\mathbf{R}|$
- D.  $2|\mathbf{R}|$
- E.  $|\mathbf{R}|^2$

- d. [1pt] What is the *largest* that  $|\mathbf{R} \bowtie \pi_{A \rightarrow B, B \rightarrow A}(\mathbf{R})|$  can be?

- A. 0
- B.  $|\mathbf{R}|$
- C.  $\frac{1}{2}|\mathbf{R}|$
- D.  $2|\mathbf{R}|$
- E.  $|\mathbf{R}|^2$

- 
- e. [1pt] The technique of *synthesis* for normalization always
- A. works only if there will be no multi-attribute keys.
  - B. achieves a 3NF schema, but it may not be dependency-preserving.
  - C. achieves a dependency-preserving, 3NF schema.
  - D. achieves a BCNF schema, but it may not be dependency-preserving.
  - E. achieves a dependency-preserving, BCNF schema.
- 
- f. [1pt] Which of the following SQL queries is illegal?
- A. `select * from T;`
  - B. `select count(*) from T;`
  - C. `select count(*) from T group by A;`
  - D. `select count(*), max (B) from T group by A;`
  - E. `select max(count(*)) from T group by A;`
- 
- g. [1pt] Which of the following SQL queries is illegal?
- A. `select A from T;`
  - B. `select A, count(*) from T;`
  - C. `select A, count(*) from T group by A;`
  - D. `select A, count(*) from T group by A, B;`
  - E. `select A, B, count(*) from T group by A, B;`
- 
- h. [1pt] Consider the schema  $\mathbf{R}(\underline{A}, B)$ ,  $\mathbf{S}(\underline{A}, \underline{D})$ , and  $\mathbf{T}(\underline{D}, B)$ . One of the following relational-algebra expressions is not like the others. That is, one of them may evaluate differently from the other four. Which one?
- A.  $\pi_B(\mathbf{R} \bowtie \mathbf{T} \bowtie \mathbf{S})$
  - B.  $\pi_B((\mathbf{R} \bowtie \mathbf{T}) \cap (\mathbf{S} \times \pi_B(\mathbf{T})))$
  - C.  $\pi_B(\mathbf{R} \bowtie \mathbf{S}) \cap \pi_B(\mathbf{S} \bowtie \mathbf{T})$
  - D.  $\pi_B(\mathbf{R} \bowtie \mathbf{S} \bowtie \mathbf{S} \bowtie \mathbf{T})$
  - E.  $\pi_B(\mathbf{R} \bowtie \mathbf{S} \bowtie \mathbf{T})$
-

- i. [1pt] Consider table **R**(A,B) for which B is of type *integer* and  $|\mathbf{R}| = n > 0$ . How many tuples will the query

```
select A from R where B <= 5 or B > 5;
```

return?

- A. 0
  - B.  $\frac{1}{2}n$
  - C.  $n$
  - D.  $n^2$
  - E. *There is not enough information to answer this.*
- 

- j. [1pt] Consider the relation **Enrol** with attributes *sid*, *cid*, *term*, and *grade* which stores academic records of students. Attribute *sid* is a student identifier and *cid* is a class—a given section of a course in a given term—identifier.

Here is a query involving **Enrol**:

```
select distinct cid
  from ( select * from Enrol E1
        where not exists
          (select *
           from Enrol E2
           where E2.cid = E1.cid
            and E2.grade > E1.grade)
        ) as V
  where grade = 7;
```

Which of the following queries must return the same result as the query above?

- I. 

```
select distinct E.cid
  from Enrol E
  where E1.grade = 7
except
select distinct E.cid
  from Enrol E
  where E1.grade > 7
```
- II. 

```
select distinct cid
  from Enrol
 group by cid
 having max(grade) = 7;
```

- A. I only.
  - B. II only.
  - C. Both I and II.
  - D. Neither I nor II.
  - E. There is not enough information available to determine this.
-

EXTRA SPACE

EXTRA SPACE

RELAX. TURN IN YOUR TEST. RETURN TO THE WILD!

## REFERENCE

*(Detach this page for convenience, if you want.)*

## Schema for the Colours Database.

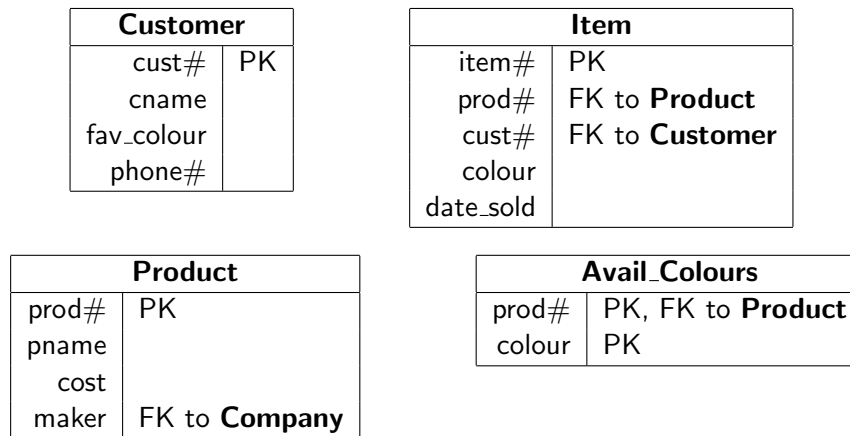


Figure 1: Colours Schema.

## Schema for the Movie Database.

**Person**(p#, name, birthdate, nationality, gender)  
**Actor**(p#, aguild#)  
 FK (p#) refs **Person**  
**Director**(p#, dguild#)  
 FK (p#) refs **Person**  
**Writer**(p#, wguild#)  
 FK (p#) refs **Person**  
**Studio**(name)  
**ScreenPlay**(title, year)  
**Authored**(title, year, writer)  
 FK (title, year) refs **ScreenPlay**  
 FK (writer) refs **Writer** (p#)  
**Movie**(title, studio, year, genre, director, length)  
 FK (studio) refs **Studio** (name)  
 FK (title, year) refs **ScreenPlay**  
 FK (director) refs **Director** (p#)  
**Cast**(title, studio, year, role, actor, minutes)  
 FK (title, studio, year) refs **Movie**  
 FK (actor) refs **Actor** (p#)  
**Affiliated**(director, studio)  
 FK (director) refs **Director** (p#)  
 FK (studio) refs **Studio** (name)

Figure 2: Movie Schema.

REFERENCE

---

**The Normal Form Definitions.**

- 1NF:** Domain of each attribute is an *elementary* type; that is, not a *set* or a *record structure*.
- 2NF:** Whenever  $\mathcal{X} \mapsto A$  is a functional dependency that holds in relation  $\mathbf{R}$  and  $A \notin \mathcal{X}$ , then either
- $A$  is *prime*, or
  - $\mathcal{X}$  is not a proper subset of any key for  $\mathbf{R}$ .
- 3NF:** Whenever  $\mathcal{X} \mapsto A$  is a functional dependency that holds in relation  $\mathbf{R}$  and  $A \notin \mathcal{X}$ , then either
- $A$  is *prime*, or
  - $\mathcal{X}$  is a key or a super-key for  $\mathbf{R}$ .
- BCNF:** Whenever  $\mathcal{X} \mapsto A$  is a functional dependency that holds in relation  $\mathbf{R}$  and  $A \notin \mathcal{X}$ , then
- $\mathcal{X}$  is a key or a super-key for  $\mathbf{R}$ .

An attribute  $A$  is called *prime* if  $A$  is in any of the candidate keys.

Figure 3: The Normal Forms.